

Which summarizer should we ship?

Our pipeline takes scanned medical records, OCRs them, extracts a structured clinical **timeline**, and generates a narrative **summary** for attorneys. Both the timeline and the summary are our core products.

The objective

We have four candidate summarizers (A, B, C, and D) behind an endpoint as **black boxes**. We need to pick one. Your task is to evaluate them and deliver two things:

1. **The selection:** Tell us which summarizer we should ship, why, and how confident you are. Your recommendation should be backed by a clear evaluation framework, not just personal preference on prose style.
2. **The application:** A working application that a reviewer can easily execute from a clean checkout to walk through the pipeline end-to-end and inspect your evaluation results.

Everything else — what you build, what you measure, what you leave out, how you present it — is yours to decide.

1 • WHAT YOU HAVE

data/<case>/ocr/*.json	the OCR text, one file per document
data/<case>/pdf/*.pdf	the same documents as scanned PDFs
data/<case>/manifest.json	what the documents are
data/<case>/summaries/*.json	real output from all four summarizers, twice each
golden/case-vance.timeline.json	a timeline for case-vance
chronos/client.py	a thin client for the two endpoints
chronos/timeline.py	the timeline format, and a checker for it
chronos/constants.py	model ids, prices, limits – the fixed facts
.env	the endpoint and your key, already filled in
CREDENTIALS.md	the same two values, and your budget

Two synthetic cases are provided (no actual PII risk). The OCR text is real pipeline output and contains typical OCR imperfections; the original scanned PDFs are included for reference. `case-vance` includes a reference timeline from our pipeline; `case-davis` does not. Pre-generated summaries under `data/<case>/summaries/` are available for you to use.

We intentionally do not provide a starter prompt or worked example. Prompt design and structured extraction are a key part of what we want to evaluate.

2 • THE TWO ENDPOINTS

Both endpoints are pre-configured in `.env` and wrapped by `chronos/client.py`. `GET /v1/models`, `/v1/cases`, and `/v1/budget` are free to call.

`POST /v1/generate` — a standard model call, and the only model you drive directly. You will use it to extract a timeline from the OCR text, and you can use it for anything else you need: evaluating, comparing, labelling, or fact-checking claims against the record.

```
{ "model": "us.anthropic.claude-sonnet-5",  
  "prompt": "...",  
  "system": "...",           // optional
```

```
"max_tokens": 4096 }

-> { "text": "...",
    "usage": { "inputTokens": 1234, "outputTokens": 567 },
    "cost_usd": 0.0129,
    "budget": { "cap_usd": 20.0, "spent_usd": 1.2, "remaining_usd": 18.8 } }
```

POST /v1/summarize/{A|B|C|D} — one of the four candidate summarizers. It has **no access to the source documents**: the summary is generated solely from the timeline you supply.

```
{ "case_id": "case-vance",
  "timeline": { "events": [
    { "date": "2023-09-04", "type": "encounter",
      "detail": "Emergency department encounter following a motor vehicle collision.",
      "source": "01_ems_and_ed p2" },
    { "date": null, "type": "diagnosis",
      "detail": "Left forearm deformity noted on arrival." }
  ]}

-> { "tool": "A", "case_id": "case-vance", "summary": "...",
    "cost_usd": 0.0243, "budget": { ... } }
```

The response includes the summary and cost, with no internal implementation details. **Note:** The summarizer identifiers (A, B, C, and D) do not mean the same thing from one case to the next. Your evaluation of any letter applies to the specific case you tested.

3 · THE TIMELINE FORMAT

Events require four fields. Extra fields are rejected to catch formatting issues early.

FIELD	
date	"2023-09-04", or null when no date is recoverable. The date the event occurred — not when it was transcribed, billed, or signed.
type	one of encounter, imaging, medication, procedure, therapy, diagnosis
detail	one sentence, required
source	optional, e.g. "04_ortho_followup p3"

Validate your timeline schema locally before submitting requests: `python -m chronos.timeline check my_timeline.json`

4 · WHAT WE WOULD LIKE BACK

A decision on which of the four summarizers we should ship, and why. State your recommendation, confidence level, and what evidence or metrics would change your mind. "Not enough evidence yet" is an acceptable conclusion if you clearly define the data you would need.

A packaged application that demonstrates your work and is reusable for future analysis. The application should accept case OCR as input so that someone running your project from a clean environment can **pick a case, generate a timeline, generate a summary, evaluate that summary**, and review the evaluation details supporting the score.

The choice of interface is yours (CLI, web app, or notebook), as long as a reviewer can execute those steps without friction.

We will spend the most time reviewing your evaluation methodology. The application is there so we can run your evaluation live and test it on new cases; it does not need polished visual design.

Submit a git repository (bundle or private link) with clear documentation. Assume reviewers will spend ten minutes reading your writeup before running the code. If you intentionally scope something out, explain your reasoning.

5 · TIME, BUDGET, TOOLS

Due	Wednesday 16 September 2026 You have it from Wednesday 9 September 2026. If you need an extension, let us know as early as possible.
------------	---

Budget	A fixed dollar budget is assigned to your API key in <code>CREDENTIALS.md</code> . Each response reports call cost and remaining balance. If you need additional budget for testing, email us with your plan — reasonable requests are always approved.
---------------	---

AI	Use it. Feel free to use Claude Code, Cursor, Copilot, or any AI tooling you prefer.
-----------	---

Two notes before you start:

1. Because model calls route through our proxy endpoint, **we can view prompt logs** (including evaluation prompts). Prompt construction and reasoning are central to how we evaluate submissions. Data is used solely for this interview.
2. The exercise materials contain deliberate inconsistencies to catch unreviewed model output. Using AI is encouraged, but failing to review generated output will hurt your score. If you notice something unexpected in the records or data, call it out in your submission.

Questions to the hiring contact on this thread.